

# Lecture 8: Inter-process Communication

Lecturer: Prof. Zichen Xu

# Outline

- Unix IPC and Synchronization
- Pipe
- Message
- Semaphore
- Shared Memory
- Signals

# Pipes and FIFOs

- Pipe: a circular buffer of fixed size written by one process and read by another
- **int pipe(int fildes[2])** : creates a pipe and returns two file descriptors, *fildes[0]* and *fildes[1]* for reading and writing
- OS enforces mutual exclusion: only one process at a time can access the pipe.
- accessed by a file descriptor, like an ordinary file
- processes sharing the pipe must have same parent in common and are unaware of each other's existence
- Unlike pipes, a FIFO has a name associated with it, allowing unrelated processes to access a single FIFO

# Pipe Example

```
main()
{
    int n;
    int pipefd[2];
    char buff[100];
    if (pipe(pipefd) < 0)    // create a pipe
        perror("pipe error");
    printf("read fd = %d, writefd = %d\n", pipefd[0], pipefd[1]);
    if (write(pipefd[1], "hello world\n", 12) != 12) // write to pipe
        perror("write error");
    if ((n=read(pipefd[0], buff, sizeof(buff))) <= 0) //read from pipe
        perror("read error");
    write(1, buff, n); /* write to stdout */
    close(pipefd[0]);
    close(pipefd[1]);
    exit(0); }
}
```

**Result:**  
read fd = 3, writefd = 4  
hello world

# Messages

- Processes read and write messages to arbitrary *message queues* (like mailboxes)
- System calls:
  - *int msgget(key\_t key, int flag)* : Creates or accesses a message queue, returns message queue identifier
  - *int msgsnd(int msqid, const void \*msgp, size\_t msgsz, int flag)* : Puts a message in the queue
  - *int msgrcv(int msqid, void \*msgp, size\_t msgsz, long msgtype, int msgflg)* : Receives a message and stores it to *msgp*
  - *msgtype*: Messages can be typed and each type defines a communication channel
  - *int msgctl(int msqid, int cmd, struct msqid\_ds \*buf)* : provides a variety of control operations on a message queue (e.g. remove)
- Process is blocked when:
  - trying to read from an empty queue
  - trying to send to a full queue

# System V IPC

- System V IPC was first introduced in SVR2, but is available now in most versions of unix
- Message Queues represent linked lists of messages, which can be written to and read from
- Shared memory allows two or more processes to share a region of memory, so that they may each read from and write to that memory region
- Semaphores synchronize access to shared resources by providing synchronized access among multiple processes trying to access those critical resources.

# Message Queues

- A Message Queue is a linked list of message structures stored inside the kernel's memory space and accessible by multiple processes
- Synchronization is provided automatically by the kernel
- New messages are added at the end of the queue
- Each message structure has a long *message type*
- Messages may be obtained from the queue either in a FIFO manner (default) or by requesting a specific *type* of message (based on *message type*)

# Message Queue Limits

- Each message queue is limited in terms of both the maximum number of messages it can contain and the maximum number of bytes it may contain
- New messages cannot be added if *either* limit is hit (new writes will normally block)
- On linux, these limits are defined as (in /usr/include/linux/msg.h):
  - MSGMAX            8192   /\*total number of messages \*/
  - MSBMNB            16384 /\* max bytes in a queue \*/



# Message Structs

- Each message structure must start with a long message type:

```
struct mymsg {  
    long msg_type;  
    char mytext[512]; /* rest of message */  
    int somethingelse;  
    float dollarval;  
};
```

# Obtaining a Message Queue

```
#include <sys/types.h>
```

```
#include <sys/ipc.h>
```

```
#include <sys/msg.h>
```

```
int msgget(key_t key, int msgflg);
```

- The key parameter is either a non-zero identifier for the queue to be created or the value IPC\_PRIVATE, which guarantees that a new queue is created.
- The msgflg parameter is the read-write permissions for the queue OR'd with one of two flags:
  - IPC\_CREAT will create a new queue or return an existing one
  - IPC\_EXCL added will force the creation of a new queue, or return an error

# Writing to a Message Queue

```
int msgsnd(int msqid, const void * msg_ptr, size_t msg_size, int msgflags);
```

- msqid is the id returned from the msgget call
- msg\_ptr is a pointer to the message structure
- msg\_size is the size of that structure
- msgflags defines what happens when no message of the appropriate type is waiting, and can be set to the following:
  - IPC\_NOWAIT (non-blocking, return -1 immediately if queue is empty)\_

# Reading from a Message Queue

```
int msgrcv(int msqid, const void * msg_ptr, size_t msg_size, long msgtype, int msgflags);
```

- msqid is the id returned from the msgget call
- msg\_ptr is a pointer to the message structure
- msg\_size is the size of that structure
- msgtype is set to:
  - = 0 first message available in FIFO stack
  - > 0 first message on queue whose type equals type
  - < 0 first message on queue whose type is the lowest value less than or equal to the absolute value of msgtype
- msgflags defines what happens when no message of the appropriate type is waiting, and can be set to the following:
  - IPC\_NOWAIT (non-blocking, return -1 immediately if queue is empty)

# Message Queue Control

```
struct msqid_ds {  
    ...                               /* pointers to first and last messages on queue */  
    __time_t msg_stime;               /* time of last msgsnd command */  
    __time_t msg_rtime;               /* time of last msgrcv command */  
    ...  
    unsigned short int __msg_cbytes; /* current number of bytes on queue */  
    msgqnum_t msg_qnum;               /* number of messages currently on queue */  
    msglen_t msg_qbytes;              /* max number of bytes allowed on queue */  
    ...                               /* pids of last msgsnd() and msgrcv() */  
};
```

- `int msgctl(int msqid, int cmd, struct msqid_ds * buf);`
- cmd can be one of:
  - `IPC_RMID` destroy the queue specified by msqid
  - `IPC_SET` set the uid, gid, mode, and qbytes for the queue
  - `IPC_STAT` get the current msqid\_ds struct for the queue

# Shared Memory

- Processes can share the same segment of memory directly when it is mapped into the address space of each sharing process
- Faster communication
- System calls:
  - *int shmget(key\_t key, size\_t size, int shmflg)* : creates a new region of shared memory or returns an existing one
  - *void \*shmat(int shmid, const void \*shmaddr, int shmflg)* : attaches a shared memory region to the virtual address space of the process
  - *int shmdt(char \*shmaddr)*:detaches a shared region
- Mutual exclusion must be provided by processes using the shared memory

# Shared Memory

- Normally, the Unix kernel prohibits one process from accessing (reading, writing) memory belonging to another process
- Sometimes, however, this restriction is inconvenient
- At such times, System V IPC Shared Memory can be created to specifically allow one process to read and/or write to memory created by another process

# Advantages of Shared Memory

- Random Access
  - you can update a small piece in the middle of a data structure, rather than the entire structure
- Efficiency
  - unlike message queues and pipes, which copy data from the process *into* memory within the kernel, shared memory is directly accessed
  - Shared memory resides in the user process memory, and is then shared among other processes



# Disadvantages of Shared Memory

- No automatic synchronization as in pipes or message queues (you have to provide any synchronization). Synchronize with *semaphores* or signals.
- You must remember that pointers are only valid within a given process. Thus, pointer offsets cannot be assumed to be valid across inter-process boundaries. This complicates the sharing of linked lists or binary trees.

# Creating Shared Memory

```
int shmget(key_t key, size_t size, int shmflg);
```

- key is either a number or the constant IPC\_PRIVATE (man ftok)
- a shmid is returned
- key\_t ftok(const char \* path, int id) will return a key value for IPC usage
- size is the size of the shared memory data
- shmflg is a rights mask (0666) OR'd with one of the following:
  - IPC\_CREAT                      will create or attach
  - IPC\_EXCL                      creates new or it will error if it exists

# Attaching to Shared Memory

- After obtaining a shmid from shmget(), you need to *attach* or map the shared memory segment to your data reference:

`void * shmat(int shmid, void * shmaddr, int shmflg)`

- shmid is the id returned from shmget()
- shmaddr is the shared memory segment address. Set this to NULL and let the system handle it.
- shmflg is one of the following (usually 0):
  - SHM\_RDONLY sets the segment readonly
  - SHM\_RND sets page boundary access
  - SHM\_SHARE\_MMU set first available aligned address

# Shared Memory Control

```
struct shmid_ds {  
    int shm_segsz;           /* size of segment in bytes */  
    __time_t shm_atime;      /* time of last shmat command */  
    __time_t shm_dtime;      /* time of last shmdt command */  
    ...  
    unsigned short int __shm_npages; /* size of segment in pages */  
    msgqnum_t shm_nattach;    /* number of current attaches */  
    ...                       /* pids of creator and last shmop */  
};
```

- `int shmctl(int shmid, int cmd, struct shmid_ds * buf);`
- cmd can be one of:
  - `IPC_RMID` destroy the memory specified by shmid
  - `IPC_SET` set the uid, gid, and mode of the shared mem
  - `IPC_STAT` get the current shmid\_ds struct for the queue

# Semaphores

- Shared memory is not access controlled by the kernel
- This means critical sections must be protected from potential conflicts with multiple writers
- A critical section is a section of code that would prove problematic if two or more separate processes wrote to it simultaneously
- Semaphores were invented to provide such locking protection on shared memory segments

# System V Semaphores

- You can create an array of semaphores that can be controlled as a group
- Semaphores may be binary (0/1), or counting
  - 1 == unlocked (available resource)
  - 0 == locked
- Thus:
  - To unlock a semaphore, you INCREMENT it
  - To lock a semaphore, you DECREMENT it
- Spinlocks are busy waiting semaphores that constantly poll to see if they may proceed

# How Semaphores Work

- A critical section is defined
- A semaphore is created to protect it
- The first process into the critical section locks the critical section
- All subsequent processes *wait* on the semaphore, and they are added to the semaphore's "waiting list"
- When the first process is out of the critical section, it *signals* the semaphore that it is done
- The semaphore then *wakes up* one of its waiting processes to proceed into the critical section
- All waiting and signaling are done *atomically*

# How Semaphores “Don’t” Work: Deadlocks and Starvation

- When two processes (p,q) are both waiting on a semaphore, and p cannot proceed until q signals, and q cannot continue until p signals. They are both asleep, waiting. Neither can signal the other, wake the other up. This is called a *deadlock*.
  - P1 locks a which succeeds, then waits on b
  - P2 locks b which succeeds, then waits on a
- Indefinite blocking, or *starvation*, occurs when one process is constantly in a wait state, and is never signaled. This often occurs in LIFO situations.



# Semaphores

- A semaphore is a non-negative integer count and is generally used to coordinate access to resources
- System calls:
  - *int sema\_init(sema\_t \*sp, unsigned int count, int type, void \* arg): Initialize semaphores pointed to by sp to count. type can assign several different types of behavior to a semaphore*
  - *int sema\_destroy(sema\_t \*sp); destroys any state related to the semaphore pointed to by sp. The semaphore storage space is not released.*
  - *int sema\_wait(sema\_t \*sp); blocks the calling thread until the semaphore count pointed to by sp is greater than zero, and then it atomically decrements the count.*
  - *int sema\_trywait(sema\_t \*sp); atomically decrements the semaphore count pointed to by sp, if the count is greater than zero; otherwise, it returns an error.*
  - *int sema\_post(sema\_t \*sp); atomically increments the semaphore count pointed to by sp. If there are any threads blocked on the semaphore, one will be unblocked.*

# Semaphores

- Example: The customer waiting-line in a bank is analogous to the synchronization scheme of a semaphore using `sema_wait()` and `sema_trywait()`:

# Semaphores example

```
#include <errno.h>

#define TELLERS 10

sema_t tellers; /* semaphore */
int banking_hours(), deposit_withdrawal;
void *customer(), do_business(), skip_banking_today();
...

sema_init(&tellers, TELLERS, USYNC_THREAD, NULL);
/* 10 tellers available */
while(banking_hours())
pthread_create(NULL, NULL, customer, deposit_withdrawal);
...

void * customer(int deposit_withdrawal)
{
int this_customer, in_a_hurry = 50;
this_customer = rand() % 100;
```

```
if (this_customer == in_a_hurry) {
    if (sema_trywait(&tellers) != 0)
        if (errno == EAGAIN) { /* no teller available */
            skip_banking_today(this_customer);
            return;
        } /* else go immediately to available teller and
            decrement tellers */
    }
    else
        sema_wait(&tellers); /* wait for next teller, then
        proceed, and decrement tellers */

    do_business(deposit_withdrawal);
    sema_post(&tellers); /* increment tellers;
                           this_customer's teller
                           is now available */
}
```

# Signals

- Software mechanism that allows one process to notify another that some event has occurred.
- Each signal is represented by a numeric value. Ex:
  - 02, SIGINT: to interrupt a process
  - 09, SIGKILL: to terminate a process
- Each signal is maintained as a single bit in the process table entry of the receiving process: the bit is set when the corresponding signal arrives
- A signal is processed as soon as the process runs in user mode

# Conclusion

- We have started on Inter-processing Communication in Linux
- We have talked a lot on different communication functions
- We will start on Linux Programming on System Security next week
- Reading Assignment: Chapter 7, 8, 9 in your textbook

# Matrix Multiplication

$$c_{i,j} = \sum_{k=1}^n a_{i,k} b_{k,j}$$

- Multiply two  $n \times n$  matrices,  $a$  and  $b$
- One each iteration, a row of  $A$  multiplies a column of  $b$ , such that:

$$C_{p,k} = C_{p,k} + a_{p,p-1} b_{p-1,k}$$